

Fake News Detection: Four Neural Networks challenge the State of the Art

Annachiara Aiello[‡] and Veronica Pistolesi[‡]

The purpose of this work is to make a contribution to the challenge of Fake News Detection: the task of correctly classifying news as fake or real. In particular, this study compares four different neural networks trained on the pre-processed Fake News Dataset^[1] with the state of the art. All the architectures are implemented with a pre-trained word embedding method, namely GloVe or fastText. The proposed models are a Convolutional Neural Network, two types of gated Recurrent Neural Networks - Bidirectional LSTM and GRU - and a combination of Convolutional Neural Network with Bidirectional LSTM. The experiments show that Bidirectional LSTM with fastText word embedding is the model with the best results.

1 Introduction

Fake news is misleading information passed off as true, with malicious intent^[2]. Although this kind of information has always existed, its spread massively increased with the recent rise of social networks. Studies state that factors implicated in the diffusion of fake news are political polarization, post-truth politics, motivated reasoning, confirmation bias, and social media algorithms.^[3] One of the most dangerous effect of fake news is that they can reduce impact of real ones, undermining people trust in serious media coverage.

Our study aims to make a contribution to the task of binary classification of news, starting from the work of I.K. Sastrawan, I.P.A. Bayupati and D.M.S. Arsa,^[4] on the Fake News Dataset^[1]. Specifically, we implemented our own Convolutional Neural Network (CNN) and Bidirectional LSTM (Bi-LSTM), with the addition of GRU and a combination of CNN with Bi-LSTM (CNN+Bi-LSTM). The purpose of the project was also to see the effect of the two pre-trained word embedding models, namely GloVe^[5] and fastText^[6]. The paper is organised into several main sections. Section 2 explains in details the connection with related works. Section 3 is further divided into three subsections: the first one and the second one describe the dataset and the word embedding methods used, respectively; the

third one presents the proposed deep learning models. Section 4 illustrates the experiment setup, Section 5 shows all the results and discussions, and eventually Section 6 contains the conclusions of the work.

2 Related Works

So far, several methods have been explored to detect fake news. Our contribution to the research is to compare the performance of four neural networks implemented with methods developed in similar studies using the same dataset, namely Fake News Dataset^[1].

Among the various experiments, Kaliyar et al.^[7] chose a pre-trained word embedding called GloVe, which was later combined with a CNN. As a result, the accuracy, precision, recall, and F1-score were 98.36%, 99.40%, 96.88%, and 98.12%, respectively.

Sastrawan et al.^[4] compared the effect of some pre-trained word embeddings, namely Word2Vec^[8], GloVe, and fastText, with three deep learning methods (CNN, Bi-LSTM and ResNet). They obtained 98.65%, 98.64%, 98.66%, and 98.65% results for accuracy, precision, recall, and F1-score, respectively, using fastText with a Bi-LSTM.

In our work, we trained our CNN and Bi-LSTM architectures, with GloVe and fastText word embedding respectively, to compare our results with those of Kaliyar et al.^[7] and of Sastrawan et al.^[4] In addition, we implemented and tested two new models, that are GRU and a combination of CNN with Bi-LSTM, both with GloVe and fastText methods. The results are shown in Section 5.

[‡] These authors are students of the master degree in Computer Science, Artificial Intelligence curriculum, University of Pisa and they contributed equally to this work.

3 Research Method

The first part of this section describes the data pre-processing and the analysis; the second one introduces the pre-trained word embedding methods; the last one illustrates the architectures of the implemented models.

3.1 Dataset

The dataset used to train and test our proposed neural networks is the Fake News Dataset^[1], which was also used in previous studies (Section 2). The dataset consists of three files, namely `train.csv`, `test.csv` and `submit.csv`. The first file (`train.csv`) is a full training dataset with the following attributes:

- id: unique id for a news article;
- title: the title of a news article;
- author: author of the news article;
- text: the text of the article; could be incomplete;
- label: a label that marks the article as potentially unreliable (1: unreliable, 0: reliable).

The second file (`test.csv`) is a testing dataset with all the same attributes as `train.csv` without the labels. The third file (`submit.csv`) contains the labels of `test.csv`. In particular, we chose to use a version of this dataset that had already been cleaned, augmented and pre-processed, provided by the best performing research team^[4] of previous studies.

In the data cleansing process, data with no content or label were deleted to remove low-quality data from the dataset. Data augmentation process were performed applying the back-translation English-German^[9] to balance the dataset. Finally, data went through pre-processing where tokenization, case folding, characters non alphabet and stopword removal and lemmatization were carried out.

At the end of these steps, the current dataset consists of two files, namely `train.csv` and `test.csv`, in which the `id` column has been dropped and the columns named `title`, `author` and `text` have been merged into a single column named `text`. Furthermore, in the case of `test.csv`, the `label` column from the `submit.csv` file was added.

Our study begins with the data analysis of the current dataset, where we considered the `test.csv` file as a black box and chose to focus our analysis on the `train.csv` file. Table 1 shows the examined values.

Exploring the data distribution is useful since neural networks need to have the inputs with similar shape and size. Furthermore, this work uses texts as input for the models and not all the sequences have the same length in terms

train.csv	
number of unique words	123659
number of news	16646
max length news	10211
min length news	2
mean length news	426
standard deviation news	443.92
median length news	317
number of outliers	561
percentage of outliers	3.37%

Table 1 Examined values in `train.csv` file

of tokens. This is why padding the final configurations of the inputs became a necessary step. Figure 1 shows the Boxplot distribution of the news length. The standard deviation of the news length showed too much dispersion in the data. Therefore, the median length of the news (dotted black line in Figure 2), which is less sensitive to outliers (3.37%), turned out to be the most representative value for our input data. For these reasons, we chose the latter as the maximum length of the padded sequences instead of the simple mean.

Final configurations of the sets, used during the model selection phase, were obtained by splitting the `train.csv` file: 80% of the data will be used as training set and the remaining 20% as validation set (data in `test.csv` will clearly be used as test set). Table 2 shows the composition of the main sets.

Data	Fake News	Real News	Total
<code>train.csv</code>	8323	8323	16646
Training Set	6658	6658	13316
Validation Set	1665	1665	3330

Table 2 Training set, validation set and full dataset composition.

3.2 Word Embedding

When words with similar meanings are visualized, they will appear next to one another thanks to a technique called word embedding or distributed word representation^[10]. This is possible due to word embedding’s ability to capture a word’s semantic and syntactic details in a sizable corpus^[11]. Because it produces promising results, word embedding is being utilized more frequently in text analysis-based research on sentiment analysis, entity recognition, part of speech tagging, and other topics^[12]. GloVe and fastText are two popular pre-trained word embedding examples.




Fig. 1 Boxplot Distribution with outliers.




Fig. 2 Countplot Distribution of the news length.

3.2.1 GloVe

Global Vector (GloVe) [5] is a widely used pre-trained word embedding. The idea behind is that it creates a vector representation of a word factorizing the matrix of the number of word occurrences in a corpus [13].

3.2.2 FastText

fastText [6] is a pre-trained word embedding that makes use of the upgraded CBOW Word2Vec architecture. Word order and the use of more effective computational procedures contributed to the improvement [14]. Since the bag of character n-gram has a positive effect on vectors representing words that are uncommon or misspelled, fastText also includes subword information [15].

3.3 Models

As already mentioned, our study tested four different deep learning methods: Convolutional Neural Network (CNN), Bidirectional LSTM (Bi-LSTM), GRU and a combination of CNN with Bi-LSTM (CNN+Bi-LSTM).

CNNs are a specialized kind of neural network for data that has a grid-like topology [16]. Since images can be thought of as a 2-D grid of pixels, they are typically used for Computer Vision problems. Nevertheless, more recently CNNs have been applied to Natural Language Processing tasks and have proven to be very effective. The idea is

that the input are sentences represented as a matrix, where each row of the matrix is a vector (in our case, a word embedding) corresponding to one token [17].

To have a direct comparison to the model proposed by Kaliyar et al. [7], we used pre-trained word embedding GloVe. We now describe our CNN architecture, which can be seen in Figure 3.

First, an embedding layer takes as input the vocabulary size, the output dimension and the input length, which is set to the median length of all news in the training set (Table 1). The embedding layer is connected to a convolutional layer, with a Rectified Linear Unit (ReLU) activation function, followed by a max-pooling layer and a dropout layer. The choice of ReLU has several advantages: it counteracts gradient vanishing, it is cheap to compute and favors sparsity [18]. Pooling reduces the output dimensions while maintaining the most salient features [17]; finally, dropout is a powerful regularization technique [16]. The above three layers are repeated two more times up to the fully connected layer and the output layer.

Recurrent Neural Networks (RNNs) are a family of neural networks for processing sequential data [16]. Among them, Long Short-term Memory (LSTMs) are devised to solve the challenge of learning long-term dependencies. In this project, two variants of LSTM were proposed: Bidirectional LSTM (Bi-LSTM) and Gated Recurrent Units (GRU).

Bi-LSTM is basically two LSTMs positioned in opposite

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 317, 300)	35958600
conv1d (Conv1D)	(None, 315, 64)	57664
max_pooling1d (MaxPooling1D)	(None, 157, 64)	0
dropout (Dropout)	(None, 157, 64)	0
conv1d_1 (Conv1D)	(None, 155, 64)	12352
max_pooling1d_1 (MaxPooling1D)	(None, 77, 64)	0
dropout_1 (Dropout)	(None, 77, 64)	0
conv1d_2 (Conv1D)	(None, 75, 64)	12352
max_pooling1d_2 (MaxPooling1D)	(None, 37, 64)	0
dropout_2 (Dropout)	(None, 37, 64)	0
flatten (Flatten)	(None, 2368)	0
dense (Dense)	(None, 64)	151616
dense_1 (Dense)	(None, 1)	65
=====		
Total params: 36,192,649		
Trainable params: 36,192,649		
Non-trainable params: 0		

Fig. 3 CNN architecture.

directions and connected to an output layer. This provides the complete past and future context information for each point in the input layer, improving memory capability of the network^[19]. To have a direct comparison to the model proposed by Sastrawan et al.^[4], we used pre-trained word embedding fastText. Our architecture, shown in Figure 4, consists of an embedding layer, followed by the Bi-LSTM layer, a global max-pooling and a dropout layer. It finishes as usual with the fully connected layer and the output layer.

In the last two architectures, we tried both GloVe and fastText word embeddings, to find out which was best suited for each. Furthermore, we experimented batch normalization, which is an adaptive reparameterization method that reduces the problem of coordinating updates across many layers of deep models^[16].

GRUs were born to answer the question of which gates of the LSTMs are really essential^[16]. Indeed, GRUs are a simplification of LSTMs, so they have far fewer parameters and require less computational time. The architecture can be seen in Figure 5 and it has an embedding layer, the GRU layer, dropout, the fully connected layer and the output layer.

At the end, we implemented an hybrid model made up of CNN and Bi-LSTM (CNN+Bi-LSTM), each of which has the

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 317, 300)	35958600
bidirectional (Bidirectional)	(None, 317, 128)	186880
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 32)	4128
dense_1 (Dense)	(None, 1)	33
=====		
Total params: 36,149,641		
Trainable params: 36,149,641		
Non-trainable params: 0		

Fig. 4 Bi-LSTM architecture.

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 317, 300)	35958600
gru (GRU)	(None, 64)	70272
batch_normalization (Batch Normalization)	(None, 64)	256
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 128)	8320
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 36,037,577		
Trainable params: 36,037,449		
Non-trainable params: 128		

Fig. 5 GRU architecture.

same structure already described (Fig. 6).

4 Experiment Setup

The project is entirely developed on Google Colab, in order to exploit its GPU. It is organized into different .ipynb files containing data processing and analysis, model selections and retraining of the best final models. Several python libraries were used, such as Tensorflow, Keras Tuner, Numpy, Seaborn, Sklearn, Pandas and Matplotlib.

For every Deep Neural Network, a model selection was made on many hyper-parameters, to select the values giving the highest performance on the validation set. A random search was run using Keras Tuner^[20], with 60 trials and validation accuracy as objective. According to "Random Search for Hyper-parameter Optimization" by Bergstra and Bengio, *"randomly chosen trials are more efficient for hyper-parameter optimization than trials on a grid"*

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 317, 300)	35958600
conv1d (Conv1D)	(None, 315, 64)	57664
max_pooling1d (MaxPooling1D)	(None, 157, 64)	0
dropout (Dropout)	(None, 157, 64)	0
conv1d_1 (Conv1D)	(None, 155, 64)	12352
max_pooling1d_1 (MaxPooling1D)	(None, 77, 64)	0
dropout_1 (Dropout)	(None, 77, 64)	0
conv1d_2 (Conv1D)	(None, 75, 64)	12352
max_pooling1d_2 (MaxPooling1D)	(None, 37, 64)	0
dropout_2 (Dropout)	(None, 37, 64)	0
bidirectional (Bidirectional)	(None, 37, 128)	66048
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dropout_3 (Dropout)	(None, 128)	0
dense (Dense)	(None, 128)	16512
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 36,123,657		
Trainable params: 36,123,657		
Non-trainable params: 0		

Fig. 6 CNN+Bi-LSTM architecture.

and they obviously "require less computational time"^[21]. Furthermore, it has been proven that "for any distribution over a sample space with a finite maximum, the maximum of 60 random observations lies within the top 5% of the true maximum, with 95% probability"^[22].

Since we noticed that all the models performed well in a few epochs, we set the search with an upper bound of 20 epochs. Furthermore, we introduced an early stopping criterion using Keras, with monitor on validation accuracy and patience equals to 3 epochs. Every model has been optimized with Adam method, using binary cross entropy loss.

Once model selection was done, the winning model of each architecture was retrained with the same optimizer and the batch size selected by the random search. The retraining was performed on the entire `train.csv` file, that is, rejoining the training and validation sets used for model selection. At this stage, we set the upper limit of epochs to 30. The stopping criterion used was instead based on the train loss of the best model of the random search.

Table 3 shows the values of the hyper-parameters considered for all model selections.

hyper-parameter	values
filters	32, 64, 128
LSTM/GRU_units	32, 64, 128
rate_dropout	0.1, 0.2, 0.3, 0.4, 0.5
dense_units	32, 64, 128
batch_size	16, 32, 64, 128
learning_rate	0.0001, 0.0005, 0.001, 0.005, 0.01
clipnorm	1, 3, 5

Table 3 Values of hyper-parameters in random search.

5 Results and Discussions

Table 4 shows the results of the proposed and state of the art models on the test set.

Regarding the architectures we made a direct comparison of, it emerges that our Bi-LSTM with fastText outperforms Sastrawan et al. (our 98.71% of accuracy on their 98.65%). Different is the case of CNN with GloVe, where we obtained 98.22% of accuracy over 98.36% by Kaliyar et al. The latter result is consistent with the fact that the architecture we implemented for CNN is only a simplification of theirs, while Bi-LSTM is the same with different hyper-parameters and the addition of the dropout layer.


As concerns the other two models, GRU and CNN+Bi-LSTM, the comparison between the two word embeddings shows that GloVe helped both models to perform better than fastText. In particular, the experiment with the hybrid model did not show improvement over the individual CNN and Bi-LSTM themselves (98.12% and 98.17% with fastText and GloVe respectively).

Ultimately, the RNNs seem to be the most effective networks for this task, with GRU reaching third place in the ranking (98.49% of accuracy), even with a basic architecture. Note that, as might be expected, training GRU has a shorter computational time than Bi-LSTM (32-40 ms/step against 62).


Figure 7 shows the confusion matrix^[23] of each model. The light colors on the main diagonal indicate that almost all news in the test set were classified correctly. The dark color in the opposite diagonal shows that there are few false positives and negatives.

5.1 Focus on Best Model


From the work emerges that Bi-LSTM with fastText word embedding reaches the best test scores (98.71% of accuracy), with a high computational cost (62 ms/step). The accuracy and loss plots of model training are in Figure 8.




(a) CNN+Bi-LSTM with fastText, confusion matrix on test set.




(b) CNN+Bi-LSTM with fastText, confusion matrix on test set.




(c) CNN with GloVe, confusion matrix on test set



(d) GRU with fastText, confusion matrix on test set.



(e) GRU with GloVe, confusion matrix on test set.




(f) Bi-LSTM with fastText, confusion matrix on test set.


Fig. 7 Confusion matrices on test set of the proposed models.

Author	Word Embedding	Model	Accuracy	Precision	Recall	F1-score	Training time (ms/step)
Sastrawan et al.	fastText	CNN	96.92%	96.86%	96.98%	96.91%	-
Sastrawan et al.	fastText	ResNet	98.11%	98.11%	98.09%	98.1%	-
Proposed model	fastText	CNN+Bi-LSTM	98.12%	98.11%	98.13%	98.12%	36
Proposed model	GloVe	CNN+Bi-LSTM	98.17%	98.16%	98.18%	98.17%	31
Proposed model	GloVe	CNN	98.22%	98.21%	98.23%	98.22%	25
Proposed model	fastText	GRU	98.24%	98.26%	98.23%	98.24%	32
Kaliyar et al.	GloVe	CNN	98.36%	99.4%	96.88%	98.12%	-
Proposed model	GloVe	GRU	98.49%	98.49%	98.49%	98.49%	40
Sastrawan et al.	fastText	Bi-LSTM	98.65%	98.64%	98.66%	98.65%	-
Proposed model	fastText	Bi-LSTM	98.71%	98.71%	98.72%	98.71%	62

Table 4 Comparison of the proposed models (highlighted in blue) with the state of the art on Fake News Dataset.



(a) Accuracy plot of Bi-LSTM with fastText on training set.



(b) Loss plot of Bi-LSTM with fastText on training set.

Fig. 8 Accuracy and loss plots of Bi-LSTM with fastText on training set.

Threshold	Accuracy	Precision	Recall	F1-score
0.50	98.71%	98.71%	98.72%	98.71%
0.80	98.54%	98.54%	98.55%	98.54%
0.90	98.44%	98.44%	98.46%	98.44%
0.95	98.36%	98.37%	98.39%	98.36%
0.99	98.12%	98.14%	98.15%	98.12%

Table 5 Test results of Bi-LSTM with fastText on different classification thresholds.

To further challenge the winning model, we decided to test it by varying the classification threshold. The results, which can be seen in Table 5, prove that even with a 0.99 threshold it still achieves good performances (98.12% of accuracy).

6 Conclusions

This study compares four neural networks with the state of the art on the Fake News Dataset^[1], in the task of Fake News Detection. With two of them, namely CNN and Bi-LSTM, it directly compares itself with the work of the best performing research teams^[4,7]; with the other two, namely GRU and CNN+Bi-LSTM, it attempts to place itself in the current ranking of results.

A pre-trained word embedding is present in all the archi-

tectures. In the case of direct comparisons, the same word embedding as in the state of the art was used, in particular Bi-LSTM was tested with fastText and CNN with GloVe. On the other side, GRU and CNN+Bi-LSTM were implemented with both fastText and GloVe to determine which was better for them.

Based on the test results, our combination of Bi-LSTM and fastText outperformed all the other models with an accuracy of 98.71%. The updated ranking of results (Table 4) also shows that GloVe performs better with both GRU and CNN+Bi-LSTM than fastText, and that RNNs seem to be more suitable models for the task.

The current fake news detection system still has much room for improvement. In the future, we would like to extend our methods by trying to expand the training dataset to cover as many topics as possible, or by using data from another language, such as Italian. Once we have obtained high-performance models, we would like to make them user-friendly by creating, for example, a web extension or a mobile app, to help people to properly inform themselves.

References

- 1 I Kadek Sastrawan, Bayupati, I Putu Agung, Arsa, and Dewa Made Sri. Fake news dataset, 2021. URL <https://data.mendeley.com/datasets/945z9xkc8d/1>.
- 2 Robert Schlesinger. Fake news in reality. *News & World Report*, April 14, 2017.
- 3 Marju Himma-Kadakas. Alternative facts and fake news entering journalistic content production cycle. *Cosmopolitan Civil Societies*, 9:25–41, July 2017.
- 4 I.P.A. Bayupati I.K. Sastrawan and D.M.S. Arsa. Detection of fake news using deep learning cnn-rnn based methods. *ICT Express*, 2021. doi:<https://doi.org/10.1016/j.icte.2021.10.003>.
- 5 Jeffrey Pennington, Richard Socher, and Christopher Manning. GloVe: Global vectors for word representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, Doha, Qatar, October 2014. Association for Computational Linguistics. doi:[10.3115/v1/D14-1162](https://doi.org/10.3115/v1/D14-1162). URL <https://aclanthology.org/D14-1162>.
- 6 Tomas Mikolov, Edouard Grave, Piotr Bojanowski, Christian Puhrsch, and Armand Joulin. Advances in pre-training distributed word representations, 2017. URL <https://arxiv.org/abs/1712.09405>.
- 7 Rohit Kumar Kaliyar, Anurag Goswami, Pratik Narang, and Soumendu Sinha. Fndnet – a deep convolutional neural network for fake news detection. *Cognitive Systems Research*, 61:32–44, 2020. ISSN 1389-0417. doi:<https://doi.org/10.1016/j.cogsys.2019.12.005>. URL <https://www.sciencedirect.com/science/article/pii/S1389041720300085>.
- 8 Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space, 2013. URL <https://arxiv.org/abs/1301.3781>.
- 9 Rico Sennrich, Barry Haddow, and Alexandra Birch. Improving neural machine translation models with monolingual data, 2015. URL <https://arxiv.org/abs/1511.06709>.
- 10 Waldemar López, Jorge Merlino, and Pablo Rodríguez-Bocca. Learning semantic information from internet domain names using word embeddings. *Engineering Applications of Artificial Intelligence*, 94:103823, 2020. ISSN 0952-1976. doi:<https://doi.org/10.1016/j.engappai.2020.103823>. URL <https://www.sciencedirect.com/science/article/pii/S0952197620301974>.
- 11 Siwei Lai, Kang Liu, Shizhu He, and Jun Zhao. How to generate a good word embedding. *IEEE Intelligent Systems*, 31(6):5–14, 2016. doi:[10.1109/MIS.2016.45](https://doi.org/10.1109/MIS.2016.45).
- 12 Abu Bakr Soliman, Kareem Eissa, and Samhaa R. El-Beltagy. Aravec: A set of arabic word embedding models for use in arabic nlp. *Procedia Computer Science*, 117:256–265, 2017. ISSN 1877-0509. doi:<https://doi.org/10.1016/j.procs.2017.10.117>. URL <https://www.sciencedirect.com/science/article/pii/S1877050917321749>. Arabic Computational Linguistics.
- 13 Marwa Naili, Anja Habacha Chaibi, and Henda Hajjami Ben Ghezala. Comparative study of word embedding methods in topic segmentation. *Procedia Computer Science*, 112:340–349, 2017. ISSN 1877-0509. doi:<https://doi.org/10.1016/j.procs.2017.08.009>. URL <https://www.sciencedirect.com/science/article/pii/S1877050917313480>. Knowledge-Based and Intelligent Information & Engineering Systems: Proceedings of the 21st International Conference, KES-20176-8 September 2017, Marseille, France.
- 14 Roger Alan Stein, Patricia A. Jaques, and João Francisco Valiati. An analysis of hierarchical text classification using word embeddings. *Information Sciences*, 471:216–232, 2019. ISSN 0020-0255. doi:<https://doi.org/10.1016/j.ins.2018.09.001>. URL <https://www.sciencedirect.com/science/article/pii/S0020025518306935>.
- 15 Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov. Enriching Word Vectors with Subword Information. *Transactions of the Association for Computational Linguistics*, 5:135–146, 06 2017. ISSN 2307-387X. doi:[10.1162/tac1_a_00051](https://doi.org/10.1162/tac1_a_00051). URL https://doi.org/10.1162/tac1_a_00051.
- 16 Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- 17 Marc Moreno Lopez and Jugal Kalita. Deep learning applied to NLP. *CoRR*, abs/1703.03091, 2017. URL <http://arxiv.org/abs/1703.03091>.
- 18 Xavier Glorot, Antoine Bordes, and Yoshua Bengio. Deep sparse rectifier neural networks. In Geoffrey Gordon, David Dunson, and Miroslav Dudík, editors, *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, volume 15 of *Proceedings of Machine Learning Research*, pages 315–323, Fort Lauderdale, FL, USA, 11–13 Apr 2011. PMLR. URL <https://proceedings.mlr.press/v15/glorot11a.html>.
- 19 Guozheng Rao, Weihang Huang, Zhiyong Feng, and Qiong Cong. Lstm with sentence representations for document-level sentiment classification. *Neu-*

- rocomputing*, 308:49–57, 2018. ISSN 0925-2312. doi:<https://doi.org/10.1016/j.neucom.2018.04.045>. URL <https://www.sciencedirect.com/science/article/pii/S092523121830479X>.
- 20 Tom O'Malley, Elie Bursztein, James Long, François Chollet, Haifeng Jin, Luca Invernizzi, et al. Kerastuner. <https://github.com/keras-team/keras-tuner>, 2019.
- 21 James Bergstra and Yoshua Bengio. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(10):281–305, 2012. URL <http://jmlr.org/papers/v13/bergstra12a.html>.
- 22 Alice Zheng, Nicole Shelby, and Ellie Volckhausen. Evaluating machine learning models. *Machine Learning in the AWS Cloud*, 2019.
- 23 Stephen V. Stehman. Selecting and interpreting measures of thematic classification accuracy. *Remote Sensing of Environment*, 62(1):77–89, 1997. ISSN 0034-4257. doi:[https://doi.org/10.1016/S0034-4257\(97\)00083-7](https://doi.org/10.1016/S0034-4257(97)00083-7). URL <https://www.sciencedirect.com/science/article/pii/S0034425797000837>.